

Who Should Own the Insights Eligibility Check?

A new feature checks whether a user is eligible before we show them financial insights. Team A owns the insights data; Team B calls Team A's service to get it. The question: who should own the eligibility rule — Team A or Team B?

⚠ **This is an anti-pattern.** Pushing a domain owner's business rule out to a calling team breaks a basic principle: the owner of the data owns the rules about that data. If we do it here, we set precedent for every future team to bypass the domain owner the same way.

Why Team A should own it

One rule, one place

Team A owns the data — Team A should own who's allowed to see it. If Team B builds its own version of the rule, every other team calling this API has to rebuild it too, and the copies will eventually disagree.

Every change becomes a group project

When the rule lives with Team A, a policy change ships once. When it's scattered across every consuming team, one change means coordinating and re-testing across all of them — slower and riskier, not faster.

Banking already solved this

This is how core banking systems work today: a bank's core ledger — not each app or branch — decides whether an account is frozen, restricted, or in good standing. Every downstream system asks the core, it never re-implements the rule itself. It's the standard for a reason.

Team B's real concern is solvable

Team B doesn't want to pay for a call that returns nothing. Fix: Team A adds a fast, cheap eligibility check so Team B skips the expensive one — without moving the rule itself out of Team A.

Recommendation

Team A owns the eligibility check, exposed as a lightweight endpoint so Team B avoids unnecessary cost. Keep the rule with the data owner — don't set precedent for bypassing it.

